Lecture 2

BASICS OF NETWORKS AND PROTOCOLS - SECURITY ISSUES

Review

- Network security is very complex
  - Many sources of threats
  - Many types of vulnerabilities
  - Some are not even “network” related, but the network provides access to Oscar
- One successful attack can lead to another!
  - Needs a lot of care and sometimes paranoia

This Class

- Consider some basics of network protocols
  - Understand some vulnerabilities through some examples
- Overview of different attacks
  - Details of a couple of other attacks will be considered next week

Example

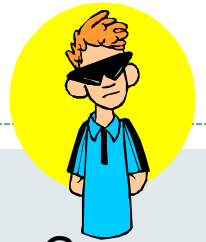

Oscar

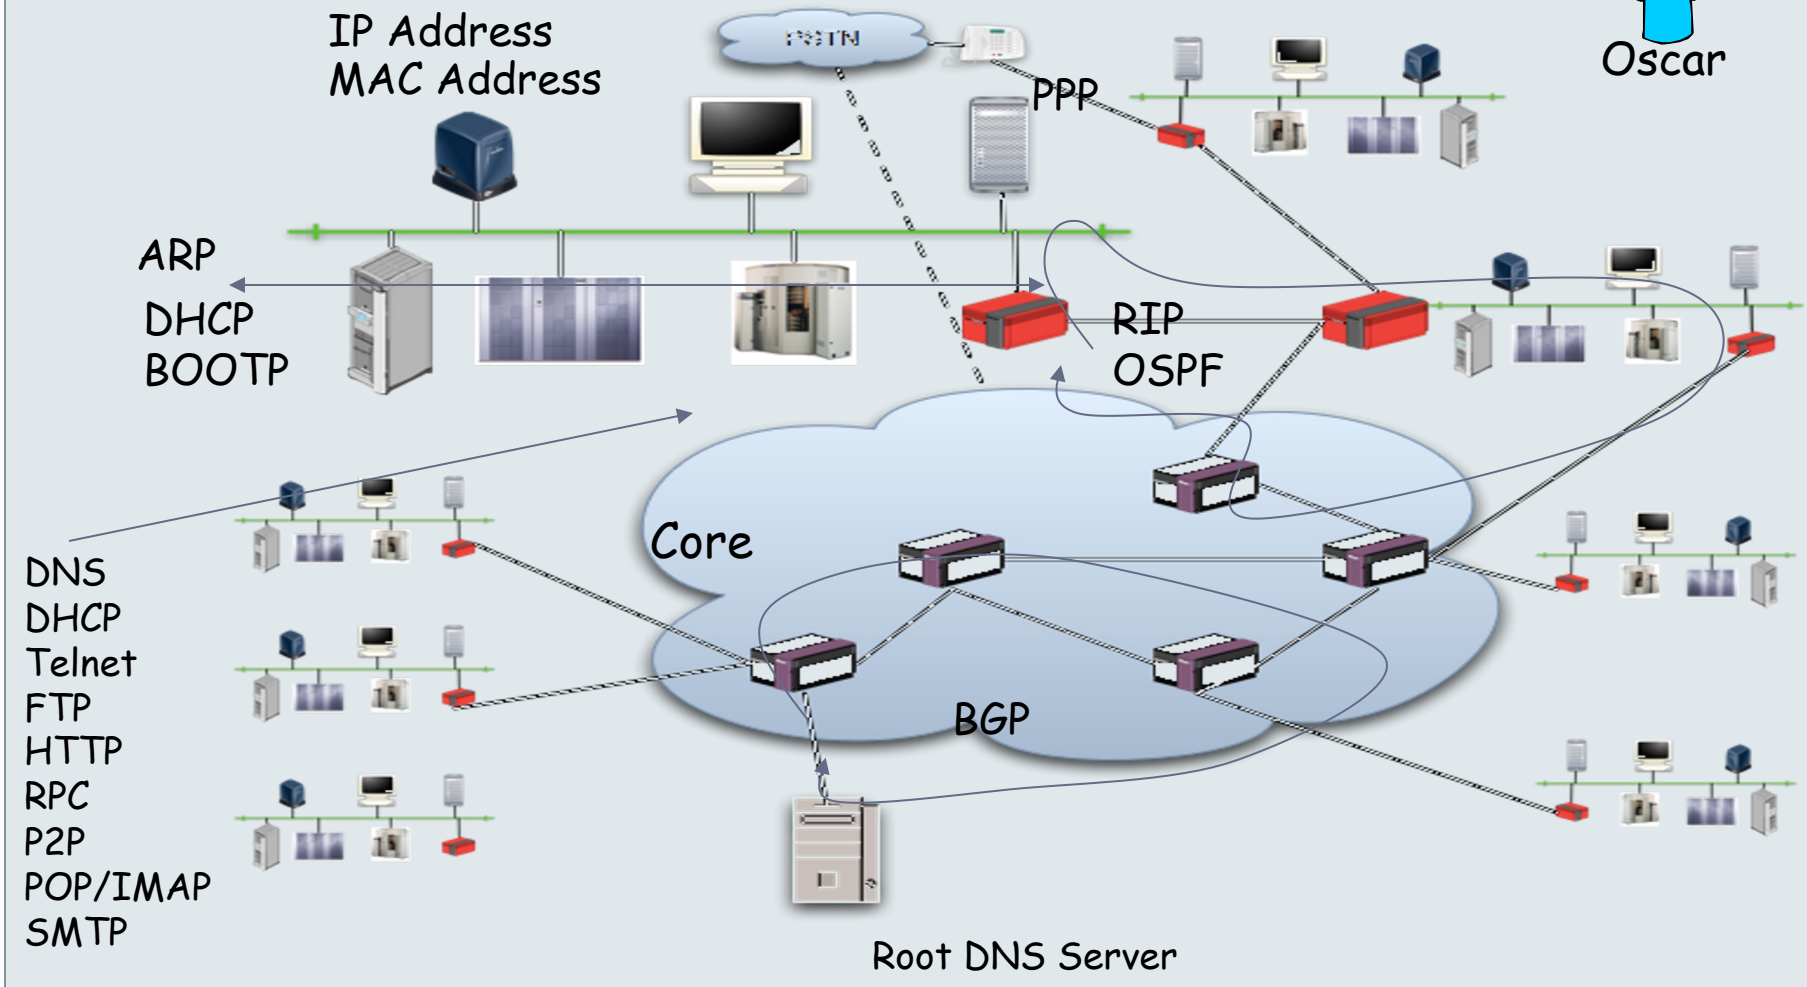

It is Complex!

- Many protocols at many layers
  - Link layer - Ethernet & 802.11 are major
  - Network layer and its “helper” protocols
    - IP, ICMP, ARP, DNS, DHCP, …
  - Transport layer - TCP and UDP are major
- Applications
  - HTTP, SMTP, FTP, Telnet, IM, RSS feeds, Other Services, Real, …

Basic Concepts

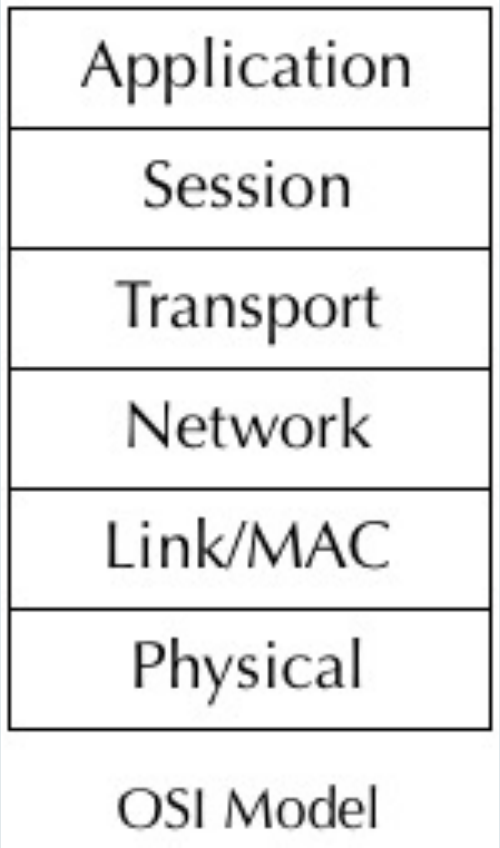

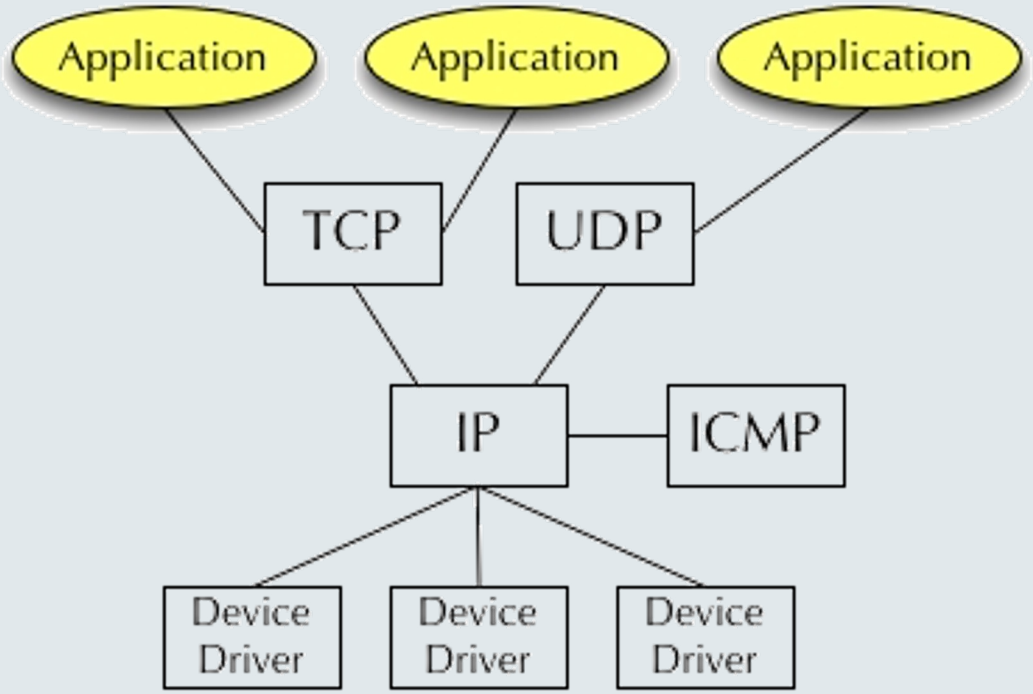

Schematic of TCP/IP Operation

Communications Across a Network

- Communicating processes typically have a client side and a server side
  - Two processes on two different hosts that communicate using *sockets*
  - A socket is like a door through which messages are sent and received
    - Interface between the application process and the transport layer
- Addressing a *process*
  - Globally unique IP address
  - Receive side port number

Processes and Sockets

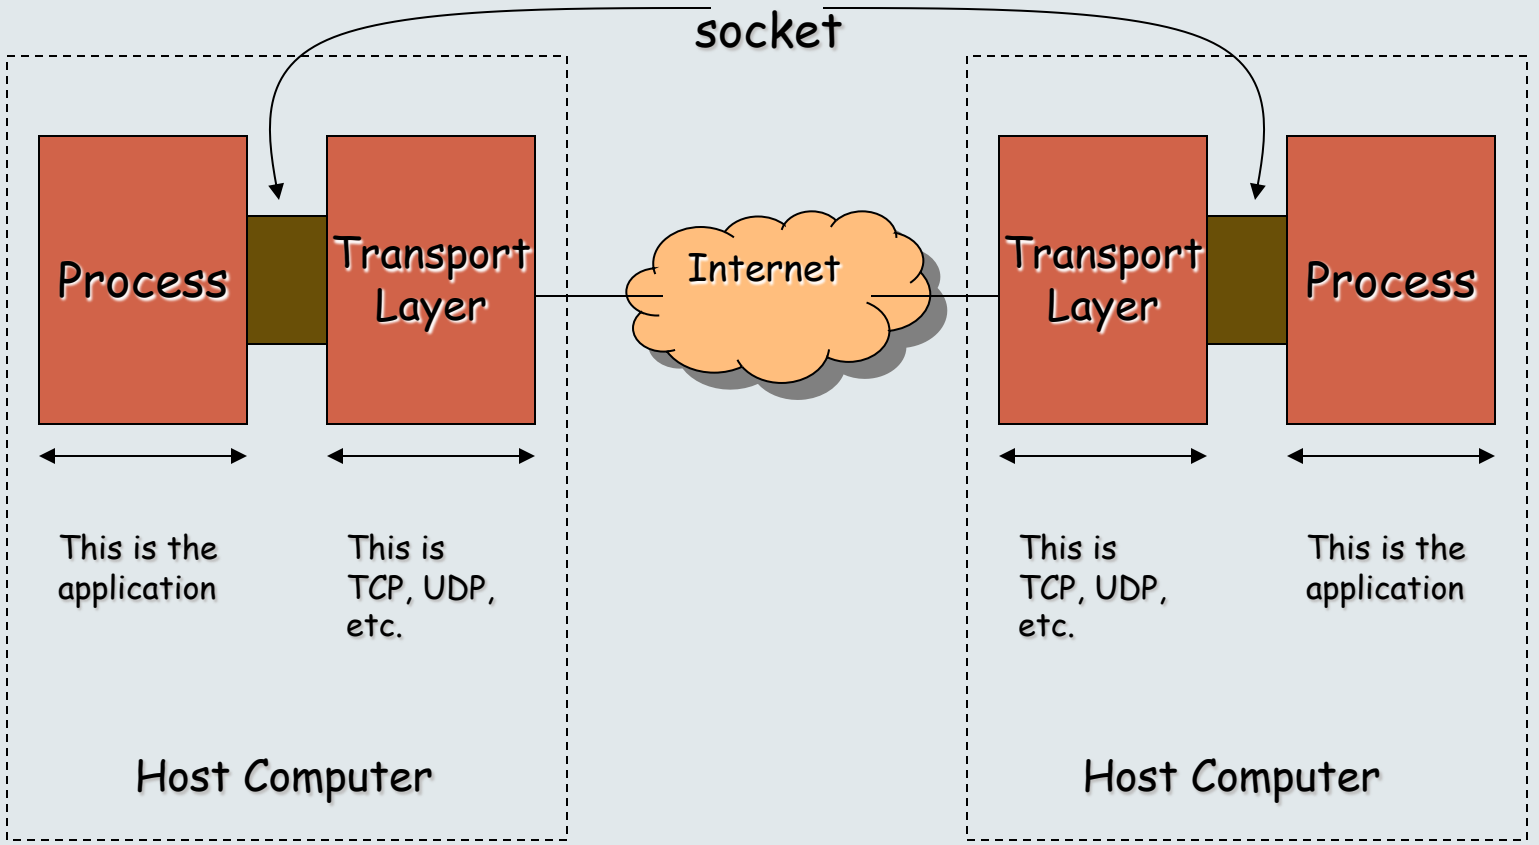

Source: Computer Networking: A top down approach by Kurose and Ross

Ports and Servers

- Client contacts the server initially for all communications
  - Server should react to the initial contact – it keeps listening to the port
    - It has an initial "socket object" to accept connections
    - It creates a new socket dedicated to a particular client after connection
  - The initial socket object is what we loosely call as an "open" port
    - It is really a half-open object
- Popular standard protocols have assigned (fixed) port numbers
  - Clients are aware of these numbers before they place a call

Port Numbers Continued

- Port numbers by convention are low numbered
  - Conventions are not always followed
  - In UNIX and UNIX-like OSs, port numbers smaller than 1024 are privileged
    - Only “root” can create these ports
    - Remote systems can trust the authenticity of these ports
- Some standard port numbers
  - Web server (http) – 80, (https) - 443;
  - DNS – 53;
  - Mail server (smtp) – 25; SSH – 22;
  - Telnet server – 23; FTP – 20 and 21;
  - POPv2 - 109, POPv3 - 110, IMAP - 143

TCP Review

- We know TCP as
  - A transport layer protocol that is carried by IP
  - A “packet” of TCP is called a segment and it is identified by a source port and a destination port
  - IP is unreliable - TCP maintains the sequence of packets in the right order and provides for acknowledgment and retransmission of lost packets
  - TCP provides flow control
    - It throttles the flow of packets if the receiver cannot handle the rate at which packets are sent
    - If a packet is dropped because of congestion, TCP will reduce the sending rate by changing the congestion window size
    - It limits the number of segments sent, but yet to be acknowledged

TCP Segment Structure

- There are six flag bits
- ACK - indicates its ACK field is valid
- RST, SYN and FIN are used for connection set up and tear down
- PSH - send data to higher layers right away
- URG - there is some urgent data

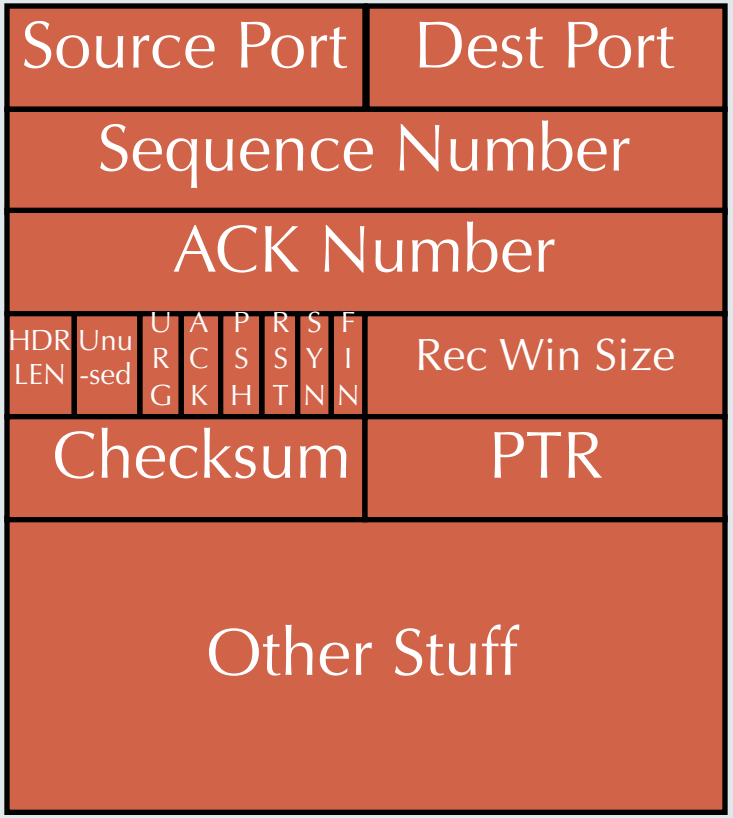

TCP Connection Management

- Client wants to initiate connection to server
  - It sends a special TCP segment to the server with the SYN bit set to 1
  - The initial sequence number is say `client_isn`
  - This is called a SYN segment
- Server receives the SYN segment
  - It allocates buffers and variables to the connection and replies
  - Reply has SYN = 1, acknowledgment number = `client_isn` +1
  - Sequence number is `server_isn`
  - This is called a SYNACK segment
- Connection is completed

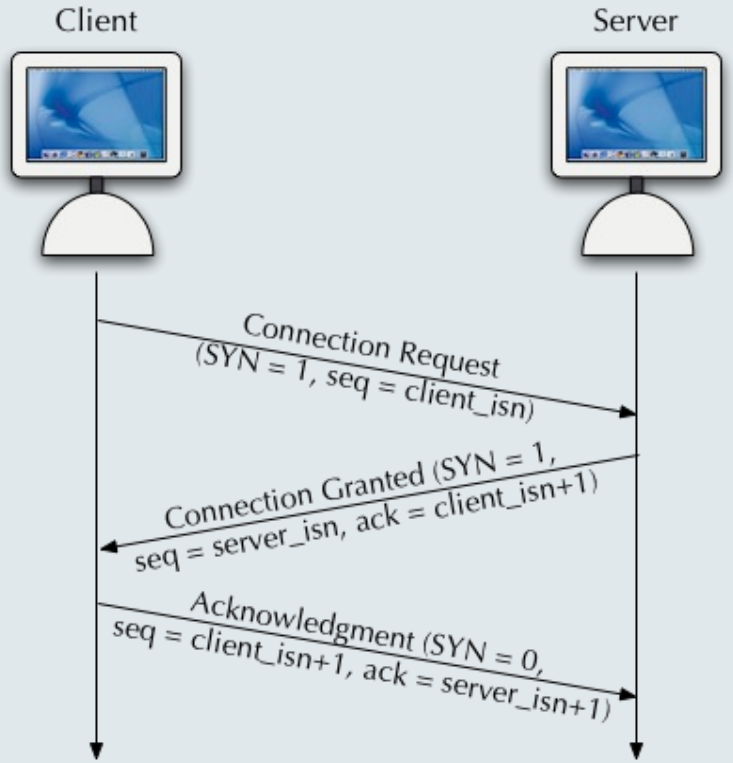

TCP States - Client and Server

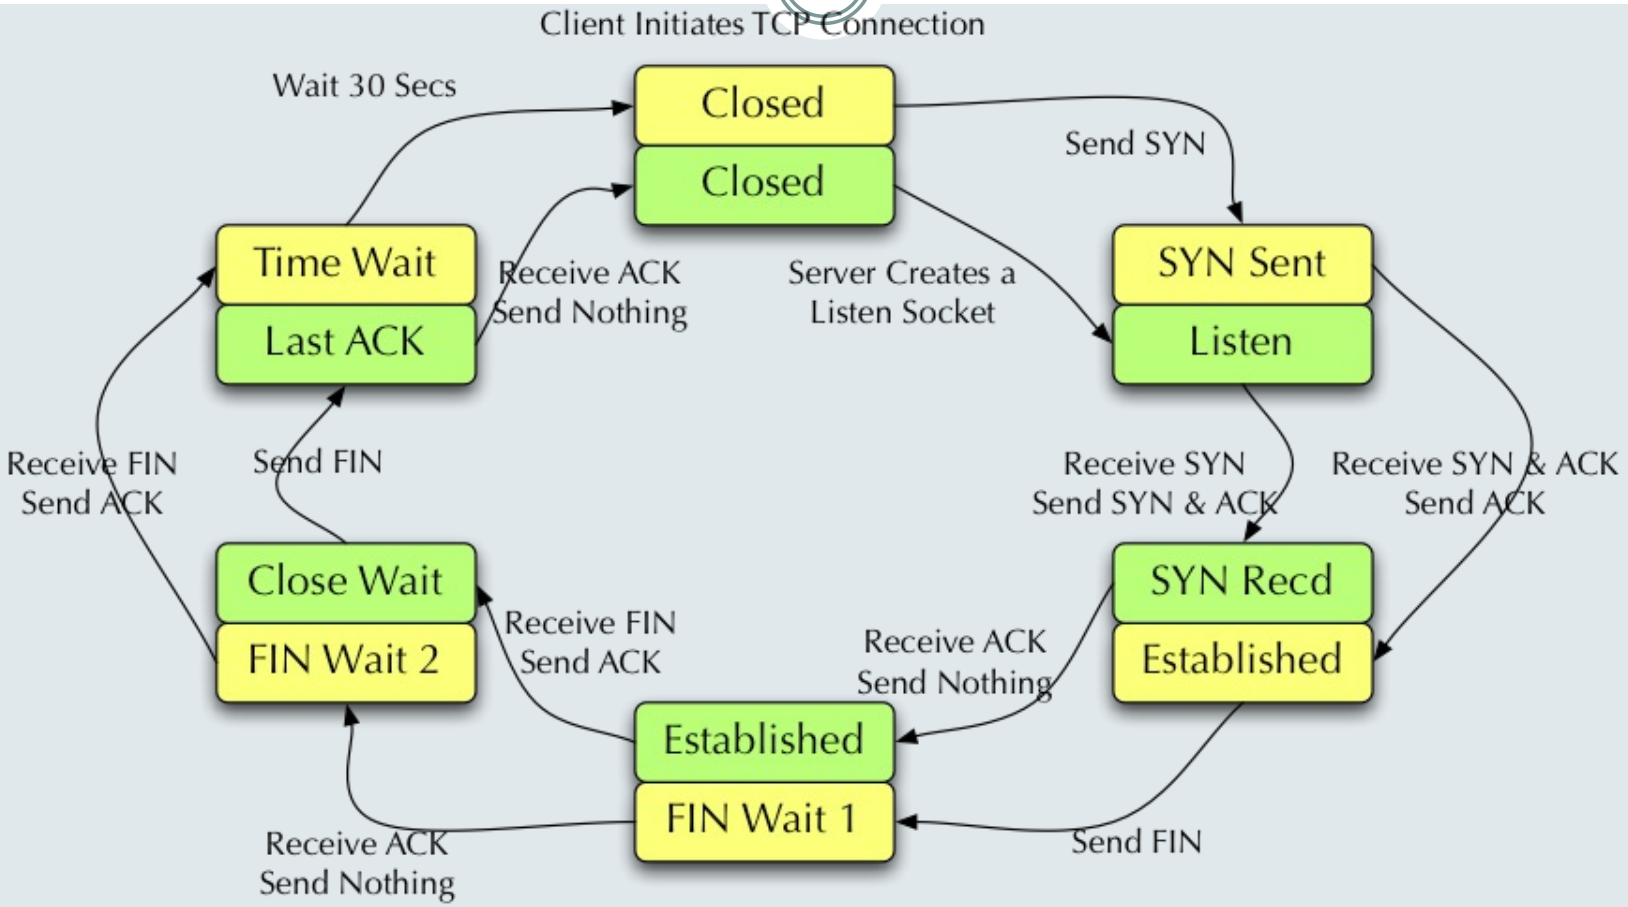

Connection Termination

- The graceful method to terminate the connection is to use the FIN field followed by ACK
  - In this case, either the client or the server will first send a TCP segment with the FIN bit set
  - The receiving host will ACK the FIN
  - This process closes *half* the connection - it has to be repeated by the receiving host
- The abrupt method of closing the TCP connection is for either the client or the server to send an RST (reset) segment
  - This aborts the TCP connection and no further communications take place between the hosts

Sequence Numbers in TCP

- Sequence and acknowledgment numbers are very important in TCP for reliable data transfer
- The sequence number of a TCP segment tells the receiver how many bytes of data has been sent
  - Example: the first TCP segment carries 1000 bytes of data and the sequence number is 235, the next TCP segment will have a sequence number 1235
- The acknowledgment number tells the recipient what is the next expected byte number
  - Example: the server receives 1000 bytes from the TCP segment with sequence number 235 - it has received bytes numbered 235 through 1234. So its sets the ack number to be 1235

ICMP

- Internet Control Message Protocol - ICMP is supposedly a very low-key protocol to answer simple requests
  - It sits below the transport layer and above the IP layer of the protocol stack
  - No port numbers of any kind - but it has types and codes in the first two bytes of the header
  - No concept of client or server - effects are mostly internal to the recipient host
  - No guarantees of delivery
- Hosts need not be listening to ICMP messages
- ICMP messages can be broadcast to hosts
- Can be a source of information leaks - e.g. host is unreachable

ICMP Codes and Types

- ICMP contains first 8 bytes of IP header that caused the ICMP response
- Ping transmits ICMP (8,0) and receives ICMP (0,0)
- Traceroute uses ICMP
  - Sends an ICMP with TTL = 1, 2, 3, 4 … to destination
  - Each router along the path detects the TTL as expired and responds with an ICMP (11,0) allowing traceroute to determine the route

Type	Code	Remark
0	0	Echo reply (ping)
3	0	Destination Network Unreachable
3	1	Destination Host Unreachable
3	3	Destination Port Unreachable
8	0	Echo request
9	0	Router advertisement
11	0	TTL Expired
12	0	IP Header Bad

Legitimate ICMP Activity

- Routers deliver “host unreachable” message
  - Common when hosts are shut down for maintenance or otherwise
  - Can be used in reconnaissance information
- Port unreachable
  - ICMP can be used to check if a UDP port is open
  - TCP ports reply with a RST/ACK flags
- Routers sometime inform you that ICMP traffic is blocked!
- Router redirect messages
  - Informs host of a more optimum router
- Need to fragment packets because MTU is exceeded
- TTL expired (time exceeded in transit)

DNS

- Domain Name System
  - Maps host names to IP addresses and vice versa
    - A tree for forward queries – What is the IP address of www.kmutnb.ac.th?
    - A tree for inverse queries – What is the host name of 136.142.116.28?
  - Common implementation is *bind*
- DNS stores so-called resource records (RRs)
  - Can reveal a lot of information about hosts and addresses

DNS Vs Typical Client-Server

- Typical client-server interaction
  - Client request connection to server
  - Server responds - handshakes take place
  - Session is initiated with interaction only between the two entities
- DNS is a bit different
  - Client issues a DNS query to the server
  - Server accepts query - may contact other DNS servers
  - Upon obtaining the information, it returns it to the client

DNS Details

- Many protocols employ DNS to translate user supplied names to IP addresses
  - DNS has to be called by http, ftp, smtp etc.
  - DNS can add delay to the communications process
- DNS is an application level protocol, but is typically not used directly by the user
- DNS queries and responses are on port 53 using UDP
  - TCP is used for zone transfers

Other DNS Services

- In addition to address mapping, DNS provides
  - Host Aliasing (e.g. www.kmutnb.ac.th can have two aliases – kmutnb.ac.th and web.kmutnb.ac.th)
  - Mail Server Aliasing (e.g. phongsakk@kmutnb.ac.th has to go to mail.kmutnb.ac.th)
  - Load Distribution (e.g. many sites use replicated web servers each running on a different end-system host)
    - DNS responds with the entire set of hosts, but rotates the order periodically

Resource Records

- Resource records (RRs) store the hostname to IP address mapping
- Each RR has four fields
  - [Name, value, type, TTL]
  - Many different types
  - TTL specifies how long the RR is valid

Name Servers

- Local Name Servers
  - Each ISP has its own name servers - all local machines contact the local name server first
  - Local translations are fast, simple and easy to implement
- Root Name Servers
  - Countable numbers worldwide (13)
  - Local servers contact the root server if they cannot resolve a name
- Authoritative Name Servers
  - Root servers direct local servers to an authoritative name server that has the information related to a host
  - Maintain authoritative data for a zone

Zone Transfers

- Zone
  - Name spaces are divided into zones based on separating “periods” in the name
  - Example: kmutnb.ac.th is a zone
- Each zone maintains primary and secondary name servers
  - Secondary servers periodically poll primary servers to obtain zone data
  - If data has changed, a zone transfer is initiated that downloads the entire database

Recursive Queries

- Local server does not know the IP address of host
  - It contacts the root server
  - The root server also does not know the IP address
    - It contacts an authoritative name server that returns the address
  - Root server returns the address to local server
- Local server forwards the IP address to requesting host
- Intermediate servers may also be used

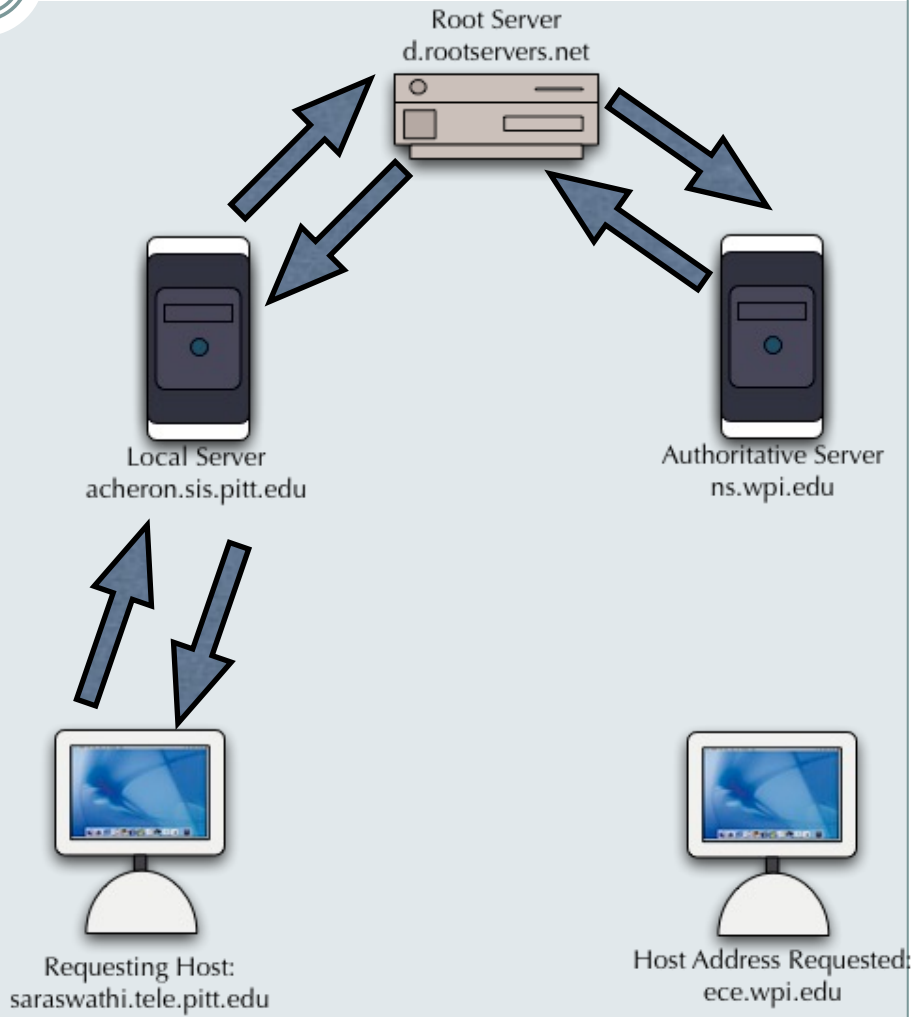

Iterative Queries

- If any server does not know the IP address, it may send the address of the next server in the list to the requesting host
- The requesting host makes direct request to the new name server
- Typically most requests are recursive, except when made to a root server
  - Query chains are a mix of iterative and recursive queries

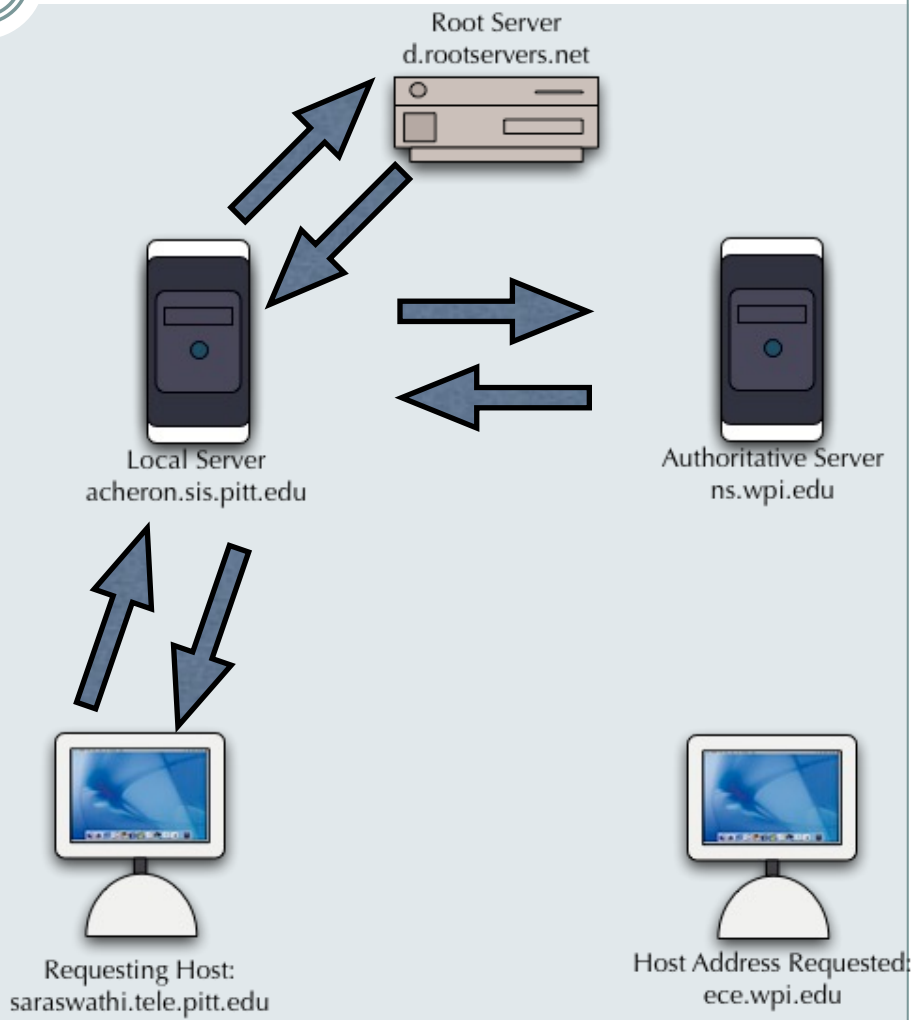

Inverse Lookup

- Inverse look-ups are performed in a slightly different way by DNS servers
- Example: Lookup 136.142.116.28
  - The query resolves 28.116.142.136.in-addr.arpa
  - Similarities between forward and inverse look-up
    - The top-level domain “arpa” has exactly one sub-domain “in-addr”
    - The host address (say 28) comes first just like forward lookups
- Inverse trees are often not current and could lead to potential security problems

DNS Software

- Berkeley Internet Name Domain (BIND)
  - Most common implementation named
  - Many versions exist (latest is 9.3.y)
- ATLAS
  - Advanced Translation Look-up And Signaling
  - Verisign's proprietary DNS software
- Microsoft has its own DNS software since Win2K
- djbdns - Free DNS software
  - See www.tinydns.org

The Security Breach Process

- Phases
  - Reconnaissance
  - Exploitation
  - Reinforcement
  - Consolidation
  - Pillage
- Network and protocol complexity and weakness aids this process

Classes of Attacks

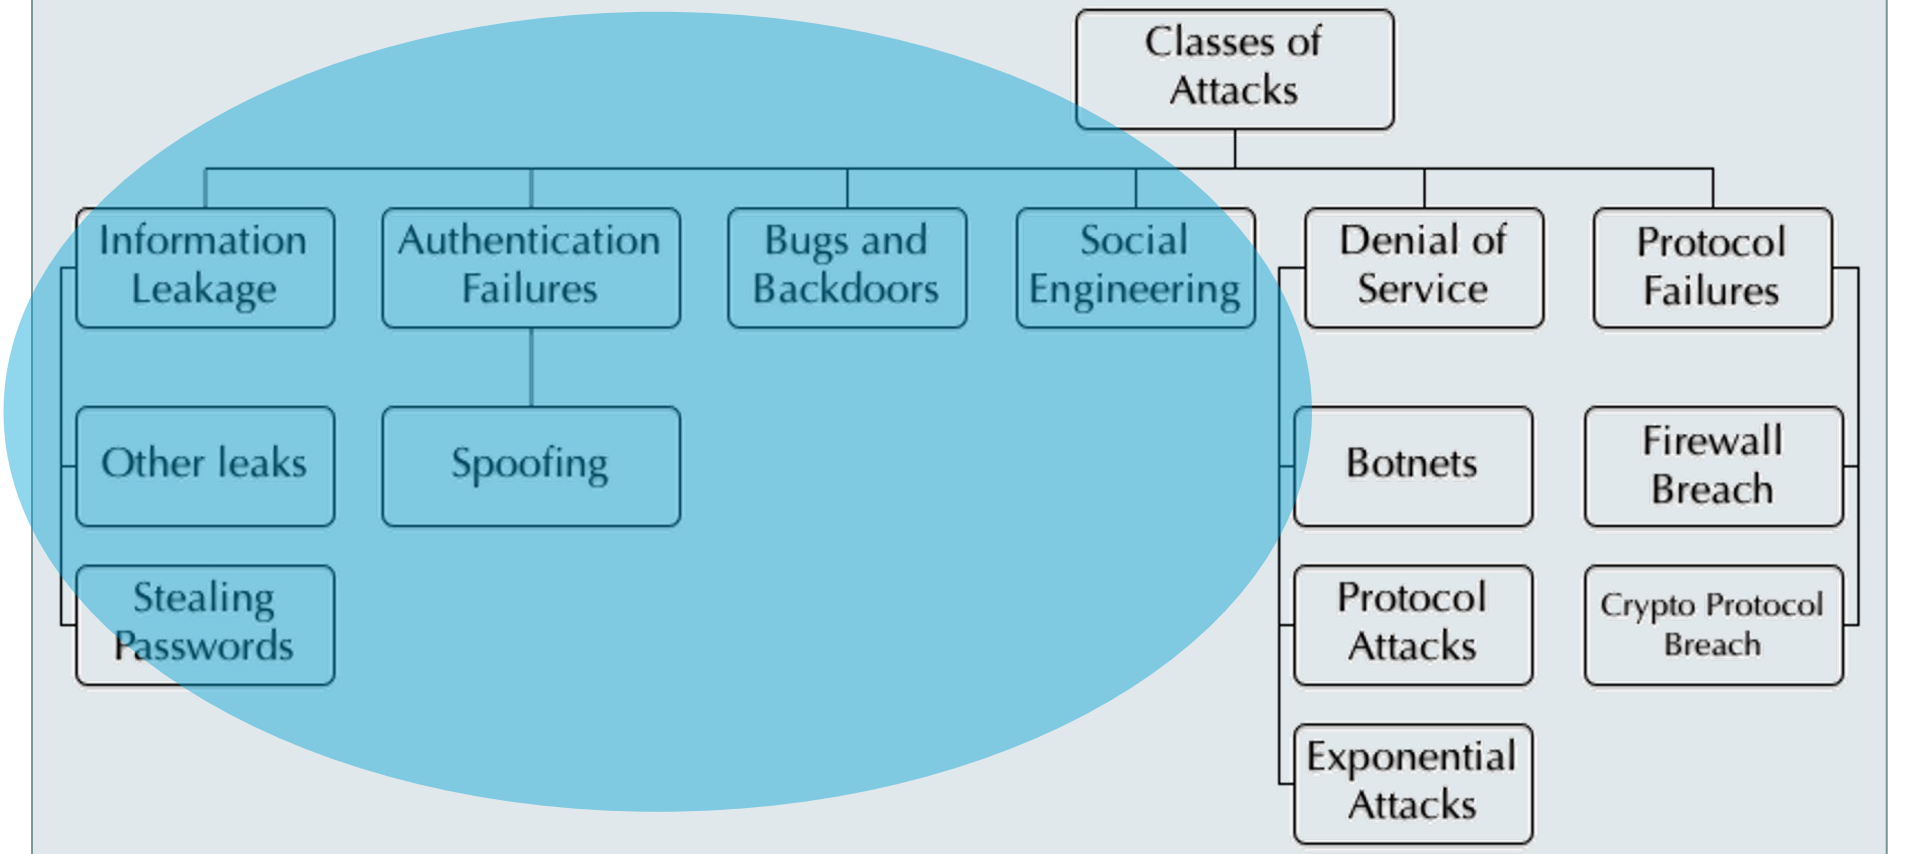

Process and Attacks - 1

- Reconnaissance
  - Makes extensive use of “Information Leakage”
  - Passive and active leakage is possible
    - Passive reconnaissance is hard to detect
    - Example: Google search
- Active reconnaissance could appear “normal”
  - Use protocols the way they are supposed to be used
- Other reconnaissance tactics are blatant
  - Port scanning, directed broadcast and so on

Process and Attacks - 2

- Exploitation, Reinforcement and Consolidation
  - Make use of stolen passwords, authentication failures, social engineering, and bugs and backdoors
- A combination of techniques can be used
  - Bugs are hard to prevent
  - Procedures and training can help prevent social engineering attacks
  - Security measures can prevent passwords from being stolen and authentication failures

Process and Attacks - 3

- Pillage
  - Good example is Denial of Service Attacks
  - We consider this later

Stealing Passwords

PACKETS FLOWING ON THE NETWORK

FILES STORED ON HOSTS ACCESSIBLE THROUGH THE NETWORK

Stealing Passwords

- Special case of information leakage
  - Larger impact on security
  - If passwords are stolen, Oscar can do a lot more than just get information
    - Oscar can move from being an outsider to a legitimate user to administrator in steps - reinforcement and consolidation
  - Harder to detect attacks because Oscar looks like a legitimate user
- Many ways in which passwords can be stolen

How passwords can be stolen

- Password in cleartext
  - Password and login are sent in cleartext by some protocols
  - Several cracker tools exist to sniff packets and get passwords
- Dictionary attacks
  - Access to the hashed password file (information leak)
  - Users typically choose a small subset of passwords – not one random password of 2^{80} choices
  - Faster to break using current technology
- Other attacks
  - Crafted Javascripts can fool users into revealing passwords
  - Other social engineering attacks

Passwords in Cleartext

- Several protocols transmit passwords in cleartext
  - Telnet
  - POP (older versions)
  - Basic authentication performed by web servers
  - SNMPv1
- Tools exist that can sniff these packets and recover passwords
  - Trivial to use and requires no knowledge of networking, protocols or programming

Dictionary Attacks

- How can Oscar get access to a hashed password?
  - Revealed /etc/passwd files
    - TFTP, SMB, NIS
    - Compromised hosts
      - Keys of ssh may also be attacked this way
  - Hashed password on the physical medium
    - POPv3
    - Digest authentication by web servers
- Password guessing and dictionary attacks
  - Given access to a password file (encrypted), Oscar tests each password to see if there is a match
  - Easy to do since the hash function is known
  - To improve the probability of success, Oscar tries common words, proper names, lowercase strings etc. – dictionary attacks
  - He can use information obtained through leakage to improve his attack!

Authentication Failures

SPOOFING ATTACKS

Authentication Failure

- Definition of sorts
  - Mechanisms to verify that the source of a request or command is legitimate fail to stop Oscar
- Common examples
  - SMTP
    - You trust that the e-mail originated from the person whose e-mail address shows up in the *From* field
    - You cannot really trust this…
  - The "r" commands
    - Commands like rlogin, rsh, etc. depend on the source address of the requesting host + assertion of username as verification of legitimacy
- Cryptographic authentication protocols are a *must*, but typically not used

IP Spoofing

- IP addresses
  - In IPv4, the address source address and destination address are both 32 bits long
  - The 32 bits are divided into two parts
    - Network portion and Host portion
  - Today people use *classless inter-domain routing* (CIDR)
    - Example: 136.142.116.28/24 means the first 24 bits are the network field
- IP address as authenticators
  - A lot of services and tools use the IP source address for authentication
  - If the IP source address is valid, trust the packet and the request!
- **You cannot rely on the validity of the source IP address except under very controlled circumstances**

ARP Spoofing

- ARP = Address Resolution Protocol
  - The link/MAC layer does not understand IP addresses
  - The NIC can only recognize MAC addresses
  - ARP is used to map the MAC address to the IP address
  - ARP packets are broadcast packets (on a LAN for example)
- If Oscar can write to the local network he can
  - Emit false ARP queries or replies
- Impact
  - Oscar can divert traffic to himself and modify data before sending it to the destination
- Notes
  - Hacker tools exist to do ARP spoofing
  - In IPv6, a “neighbor discovery” or ND protocol is used instead of ARP and can create more serious problems if spoofed

TCP and UDP

- Cannot trust privileged port numbers from TCP
  - In UNIX and UNIX-like OSs, port numbers below 1024 are “privileged”
    - Only “root” can access these numbers
  - This is meaningless for other OSs
    - Also easy to spoof the port numbers in specially crafted packets
  - We consider attacks on TCP in more detail next week
- UDP sequence numbers can be easily spoofed
  - Since there are no handshakes with UDP, it is easier to spoof UDP
  - UDP carries several services (like DNS) and can be dangerous

DNS and Authentication Failures

- DNS reverse lookup is used to authenticate the "r" commands
  - If Oscar controls the reverse lookup tree by some chance, he can falsify it
  - Inverse record will contain the name of a machine that your machine trusts
- Newer DNS lookups perform the lookup both ways to prevent such attacks
  - Cross-checking is done by the *gethostbyaddr*
  - If anomalies are detected, they should be logged

DHCP/BOOTP

- DHCP = Dynamic Host Configuration Protocol
  - Used to assign IP addresses
  - Supply information about name servers, gateways, etc.
  - Client sends a UDP broadcast request
  - Server replies with information
  - Can interface with name servers to enable mapping names to IP addresses
- Can supply a lot of information
- Logs are important for forensics
- Used only on local networks
  - Needs to know the MAC address of client
  - Reduces risks, but spoofed messages can divert traffic
  - Easier to spoof ARP and achieve the same objectives

Cookies etc.

- HTTP is stateless
  - Each HTTP request and response are treated in isolation
  - Hard for web servers to determine their state with the client they are serving
- Cookies
  - Maintain state information for servers
  - Sometimes hidden input fields or special fields in URLs are used to maintain state
  - Some web servers rely on cookies for authentication
- Cookies can be easily spoofed
  - Users can change cookies
  - Server can encrypt cookies but it is subject to other kinds of attacks (like?)
- Canned shopping carts…

Other authentication failures

- RPC and RPCBind
  - Easy to spoof userid, groupid, machine name, etc.
  - RPCbind
    - More dangerous since you can ask RPCbind to issue an indirect call to a service
  - Solutions
    - Use kerberized version of RPC
- NTP – Network Time Protocol
  - Some authentication tokens are designed to expire after a “lifetime”
    - Example: Kerberos
  - Spoofed to allow replay of authentication tokens

Bugs and Backdoors

- *Buffer overflows* are the biggest problem in creating bugs and backdoors
  - Example: finger daemon and the Internet Worm
- Protocols that have seen many bugs
  - Sendmail
  - RPC
  - NFS
  - FTP
- Another common problem is misconfiguration
  - FTP daemon
  - Anonymous FTP sites
  - Example: Java FTP client

Bugs and Backdoors - II

- Other non-obvious ways of exploiting bugs and backdoors
  - HTTP returned documents
    - May request a specific program to process them
  - Spyware, Adware, Foistware
    - No patching
  - ActiveX
    - If the code is signed, it can be trusted!
  - Browsers that allow weak ciphersuites
- Poorly written server scripts
  - Provide entry points for Oscar to insert malicious code

Social Engineering

- Read Kevin Mitnick, “The Art of Deception”
- E-mails, URLs and Javascripts
- Phishing and Pharming
- Ignorance and naiveté
- Carelessness
- FIS suggests near-paranoid behavior